

Agenda

- Introduction on Stream Processing Models [done]
- Declarative Language: Opportunities, and Design Principles [done]
- Comparison of Prominent Streaming SQL Dialects for Big Stream Processing Systems
- Conclusion

Our Focus

- Prominent Big Stream Processing Engines that offer a declarative SQL-like interface.
 - Flink,
 - Spark Structured Streaming, and
 - Kafka SQL (KSQL)

Flink SQL

- Available since version 1.3, it builds on Flink Table API (LINQ-style API)
- Uses Apache Calcite for parsing, interpreting and planning, while execution relies on FLINK Runtime.
- **Relevant concepts:** windows as group-by function, temporal tables, match-recognize (not today)

Spark Structured Streaming

- Available since Spark 2.0, it extends Dataframe and Datasets to Streaming Datasets
- SQL-like programming interface that relies on Catalyst for optimization
- **Relevant Concepts:** Complete, Append, and Update modes/

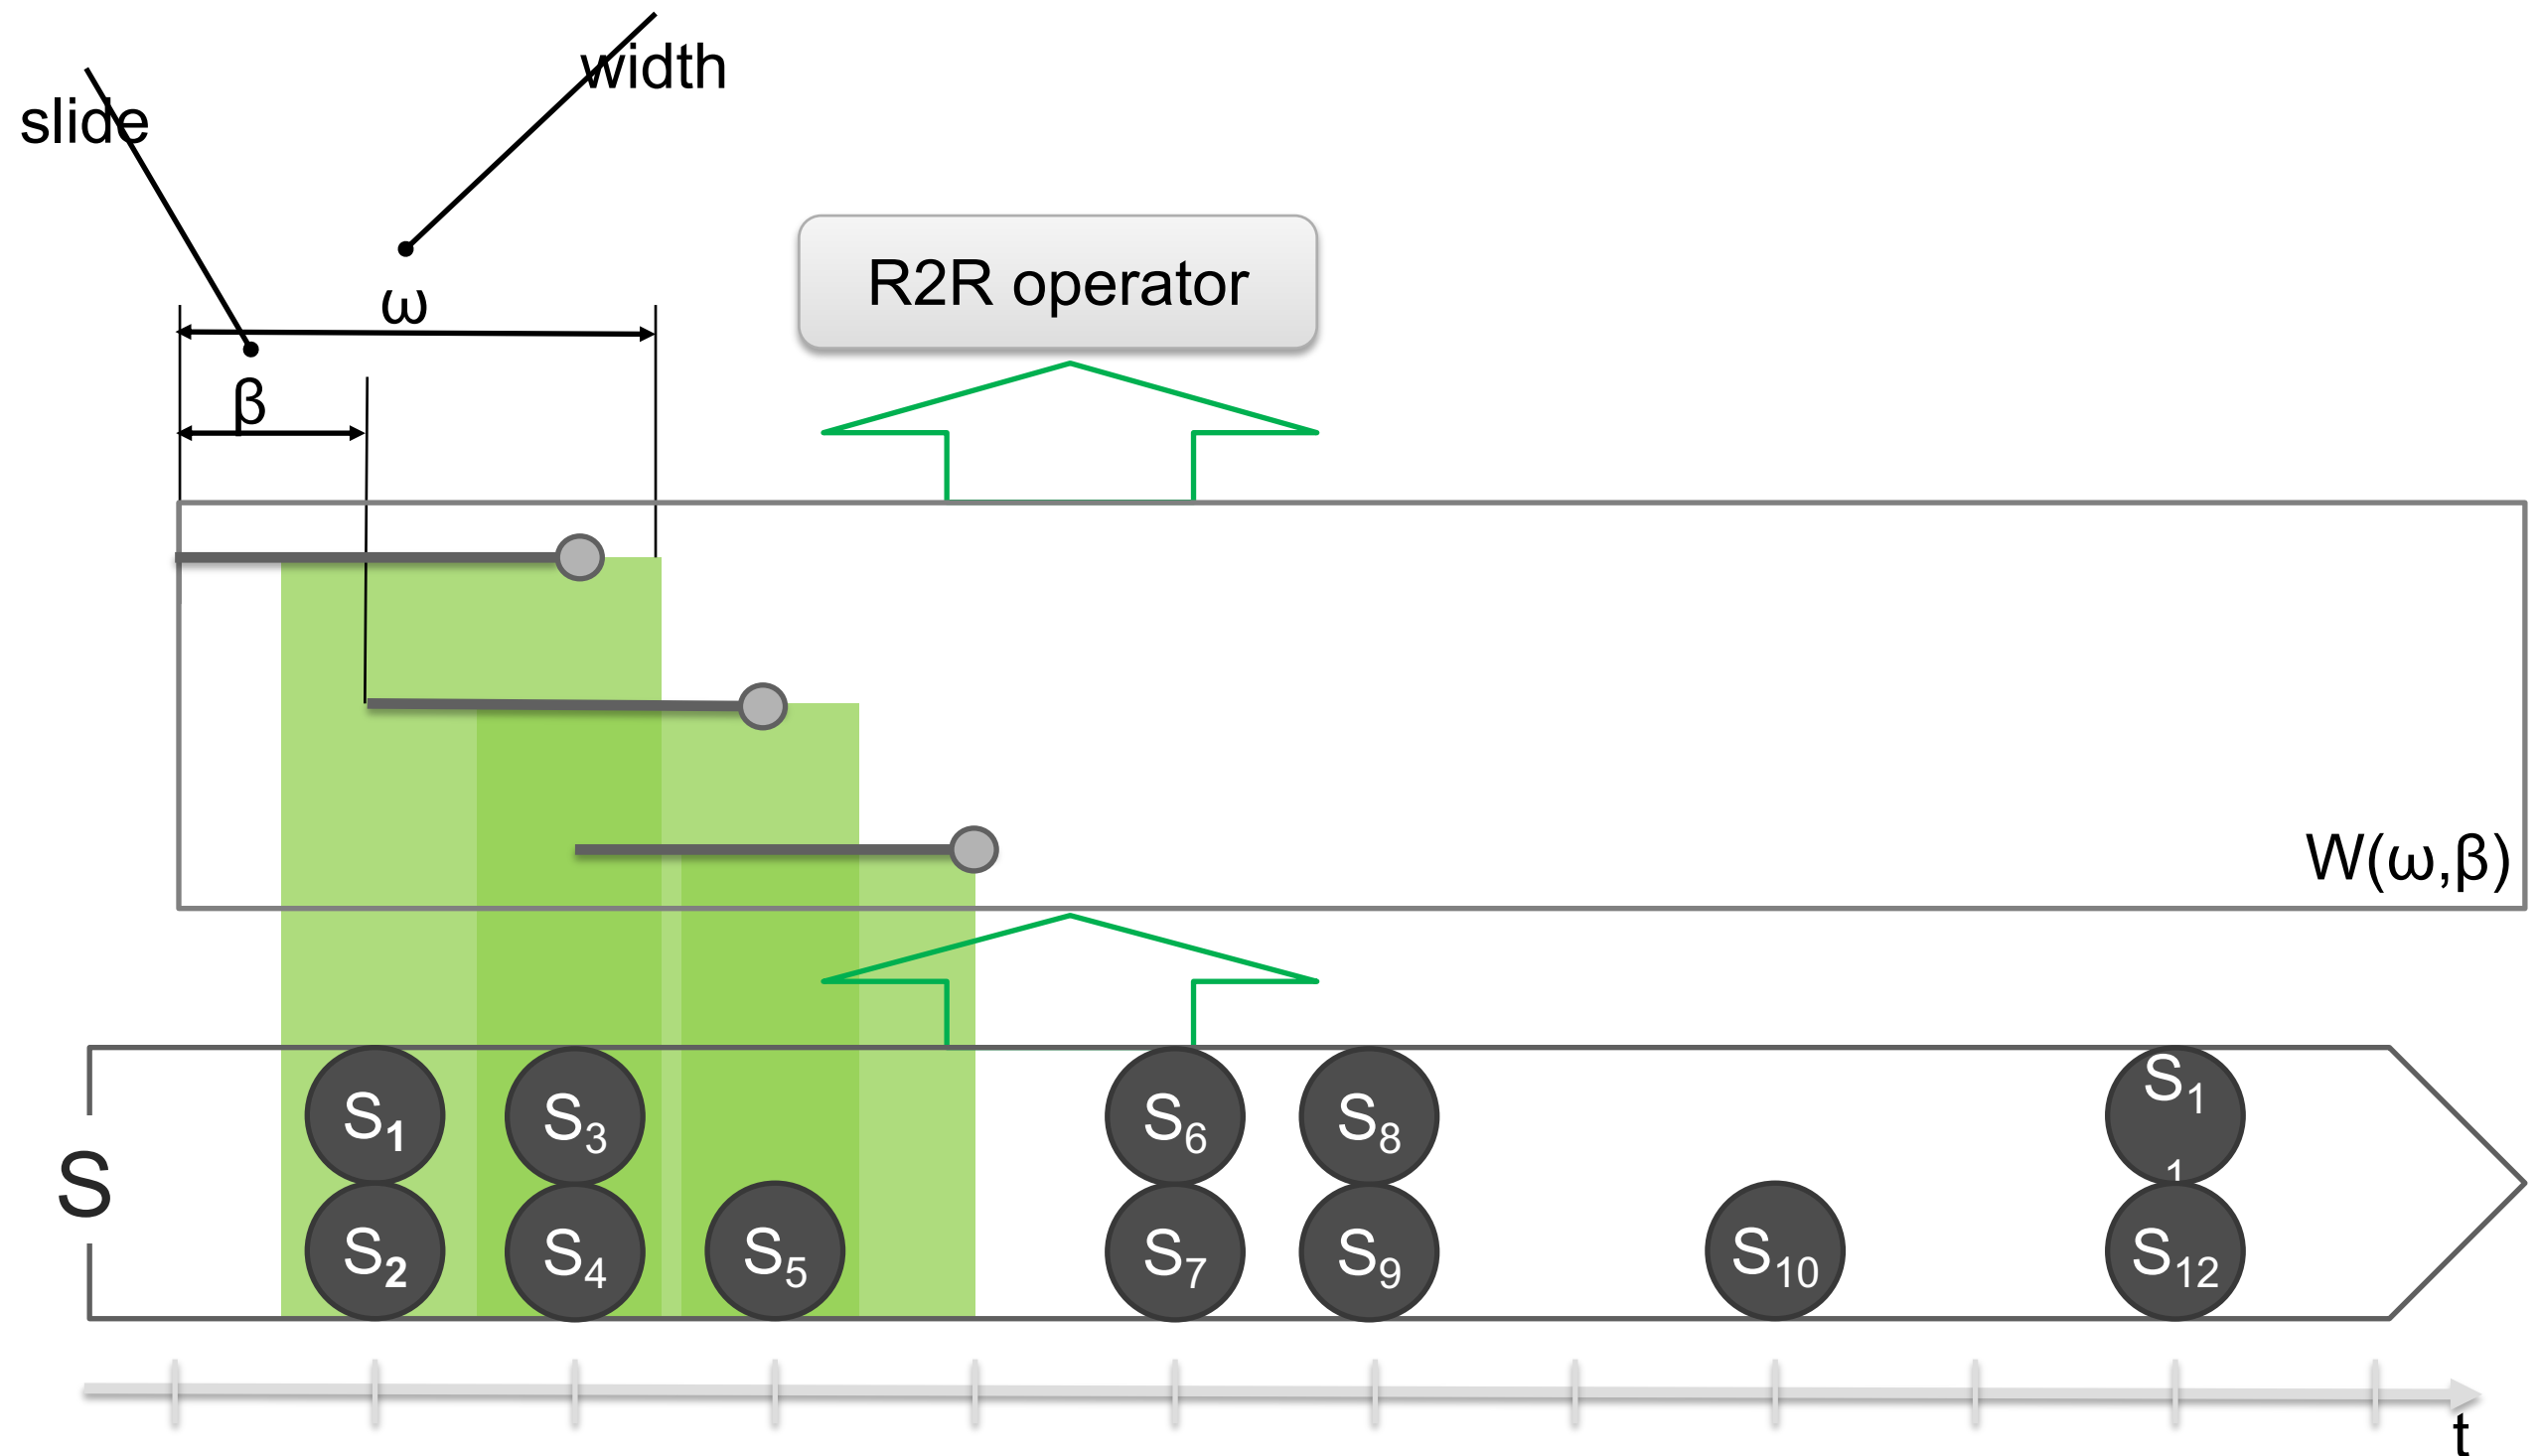
Kafka SQL (KSQL)

- Available since Kafka 1.9/2 (or confluent platform 5)
- builds directly on top of KStreams Library
- **Relevant Concepts:** simplicity is the key, relation (compacted) topic vs table/stream

Time-Window Operators & Aggregates

- Sliding Window
- Tumbling Window
- Session Window
- Aggregations: COUNT, SUM, AVG, MEAN, MAX, MIX

Sliding/Hopping Window



KSQL Hopping Window

DDL Extension

Aggregat

Window From Function

```
CREATE TABLE analysis AS  
SELECT nation, COUNT(*)  
FROM pageviews  
WINDOW HOPPING (SIZE 30 SECONDS, ADVANCE BY 10  
SECONDS)  
GROUP BY nation;
```


Results

SELECT * FROM analysis

1561375069212		Page_66	:	Window{start=1561375050000 end=-}		Page_66		1
1561375069311		Page_11	:	Window{start=1561375050000 end=-}		Page_11		1
1561375073332		Page_33	:	Window{start=1561375050000 end=-}		Page_33		1
1561375077242		Page_32	:	Window{start=1561375050000 end=-}		Page_32		1
1561375080706		Page_55	:	Window{start=1561375080000 end=-}		Page_55		1
1561375082825		Page_34	:	Window{start=1561375080000 end=-}		Page_34		1
1561375085084		Page_56	:	Window{start=1561375080000 end=-}		Page_56		1
1561375086275		Page_85	:	Window{start=1561375080000 end=-}		Page_85		1
1561375086905		Page_20	:	Window{start=1561375080000 end=-}		Page_20		1
1561375094475		Page_27	:	Window{start=1561375080000 end=-}		Page_27		1

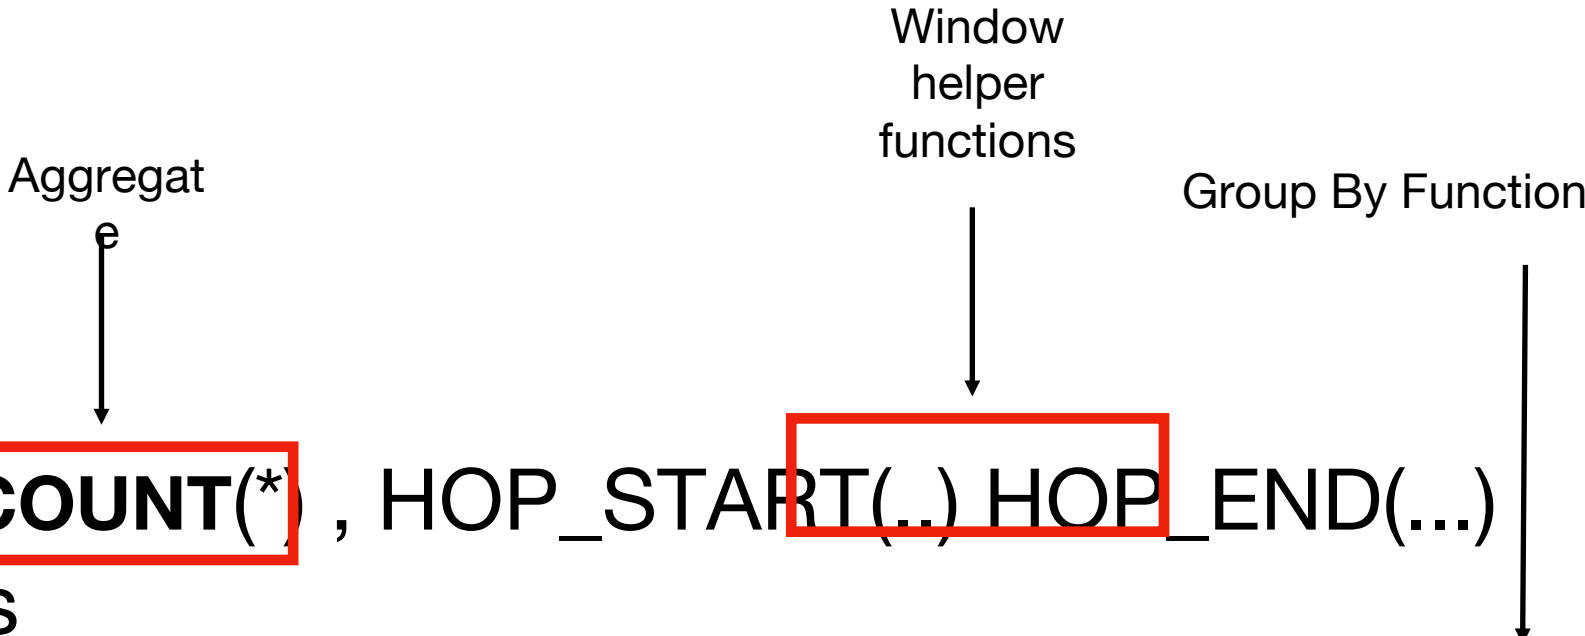
Flink SQL Hopping Window

Aggregat
e

Window
helper
functions

Group By Function

```
SELECT nation, COUNT(*) , HOP_START(..) HOP_END(...)  
FROM pageviews  
GROUP BY HOP(rowtime, INTERVAL 1H, INTERVAL 1M),  
nation
```



Results

1> (Egypt,2019-06-24 11:38:00.0,2019-06-24 11:38:01.0,1)
1> (Egypt,2019-06-24 11:39:00.0,2019-06-24 11:39:01.0,1)
1> (Egypt,2019-06-24 11:40:00.0,2019-06-24 11:40:01.0,1)
1> (Egypt,2019-06-24 11:41:00.0,2019-06-24 11:41:01.0,1)
2> (Italy,2019-06-24 11:42:00.0,2019-06-24 11:42:01.0,1)
2> (Italy,2019-06-24 11:43:00.0,2019-06-24 11:43:01.0,1)
2> (Italy,2019-06-24 11:44:00.0,2019-06-24 11:44:01.0,1)
2> (Italy,2019-06-24 11:45:00.0,2019-06-24 11:45:01.0,1)
2> (Italy,2019-06-24 11:46:00.0,2019-06-24 11:46:01.0,1)
2> (Italy,2019-06-24 11:47:00.0,2019-06-24 11:47:01.0,1)
2> (Italy,2019-06-24 11:48:00.0,2019-06-24 11:48:01.0,1)
3> (Estonia,2019-06-24 11:49:00.0,2019-06-24 11:49:01.0,1)

Spark Structured Streaming Hopping Window

```
val df = pageviews  
.groupBy(
```

```
  window($"timestamp", "1 hour", "1 minute"), $"nation").count()
```

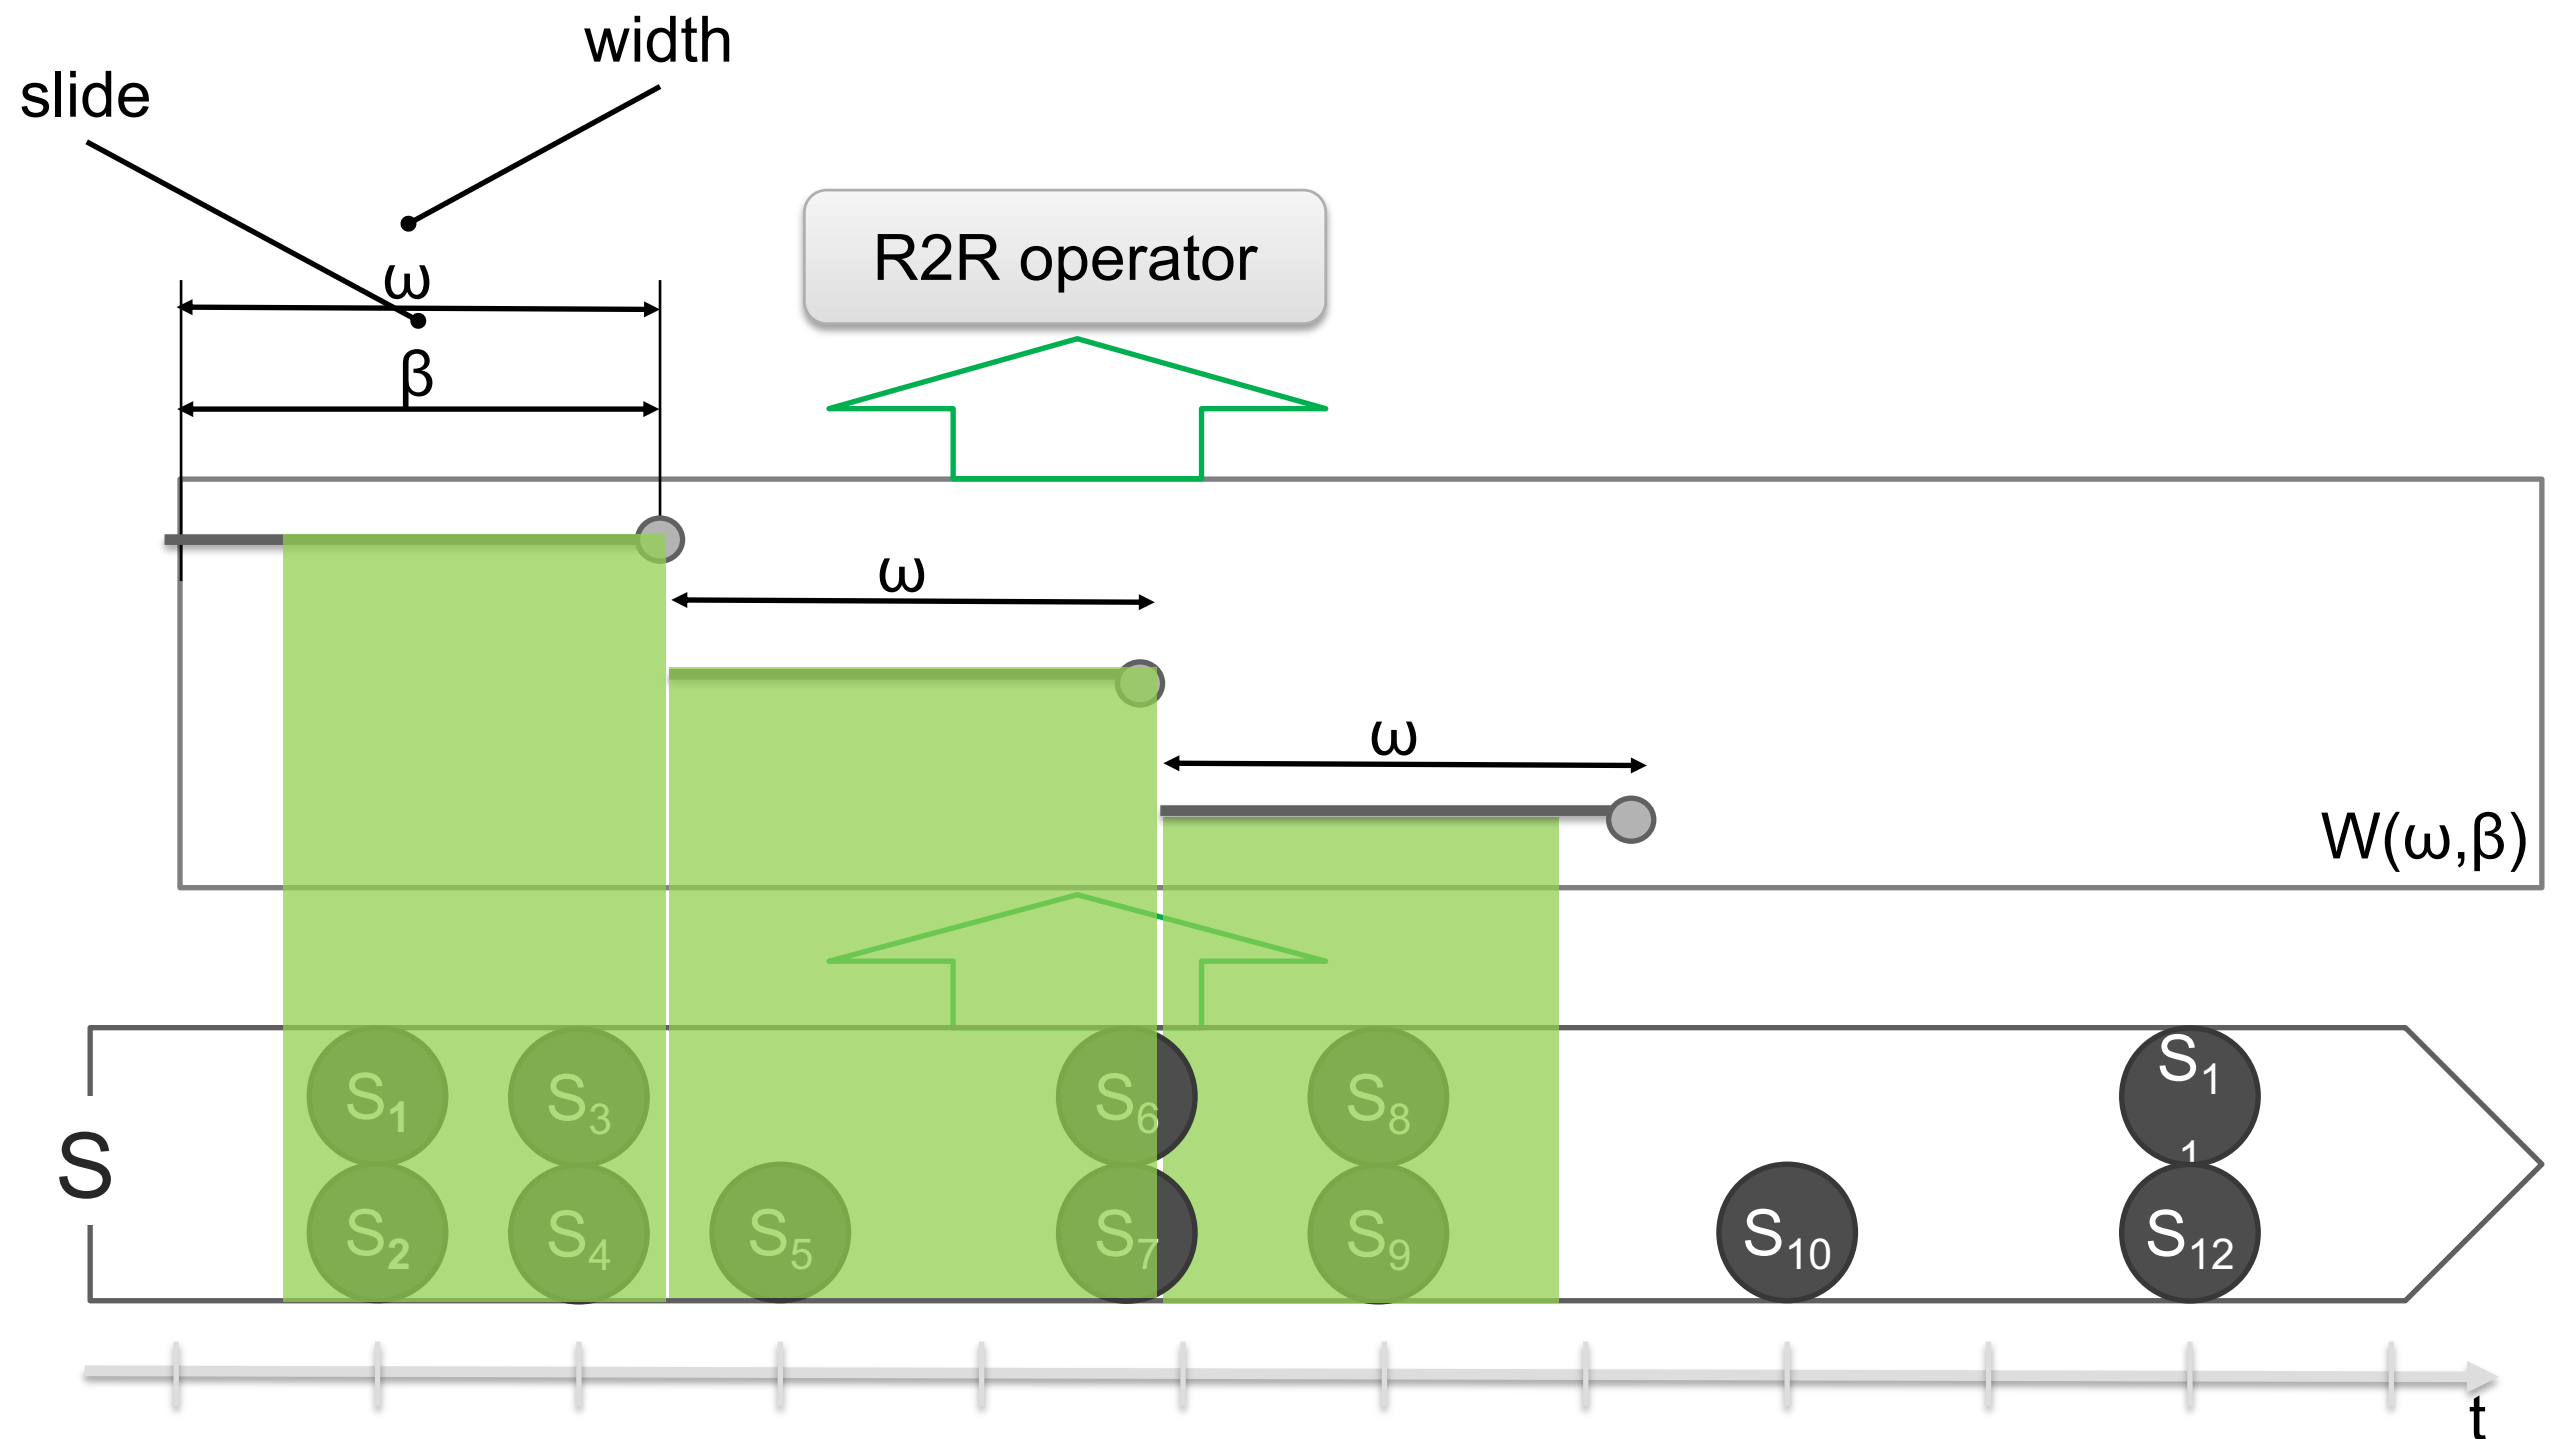
Window operator



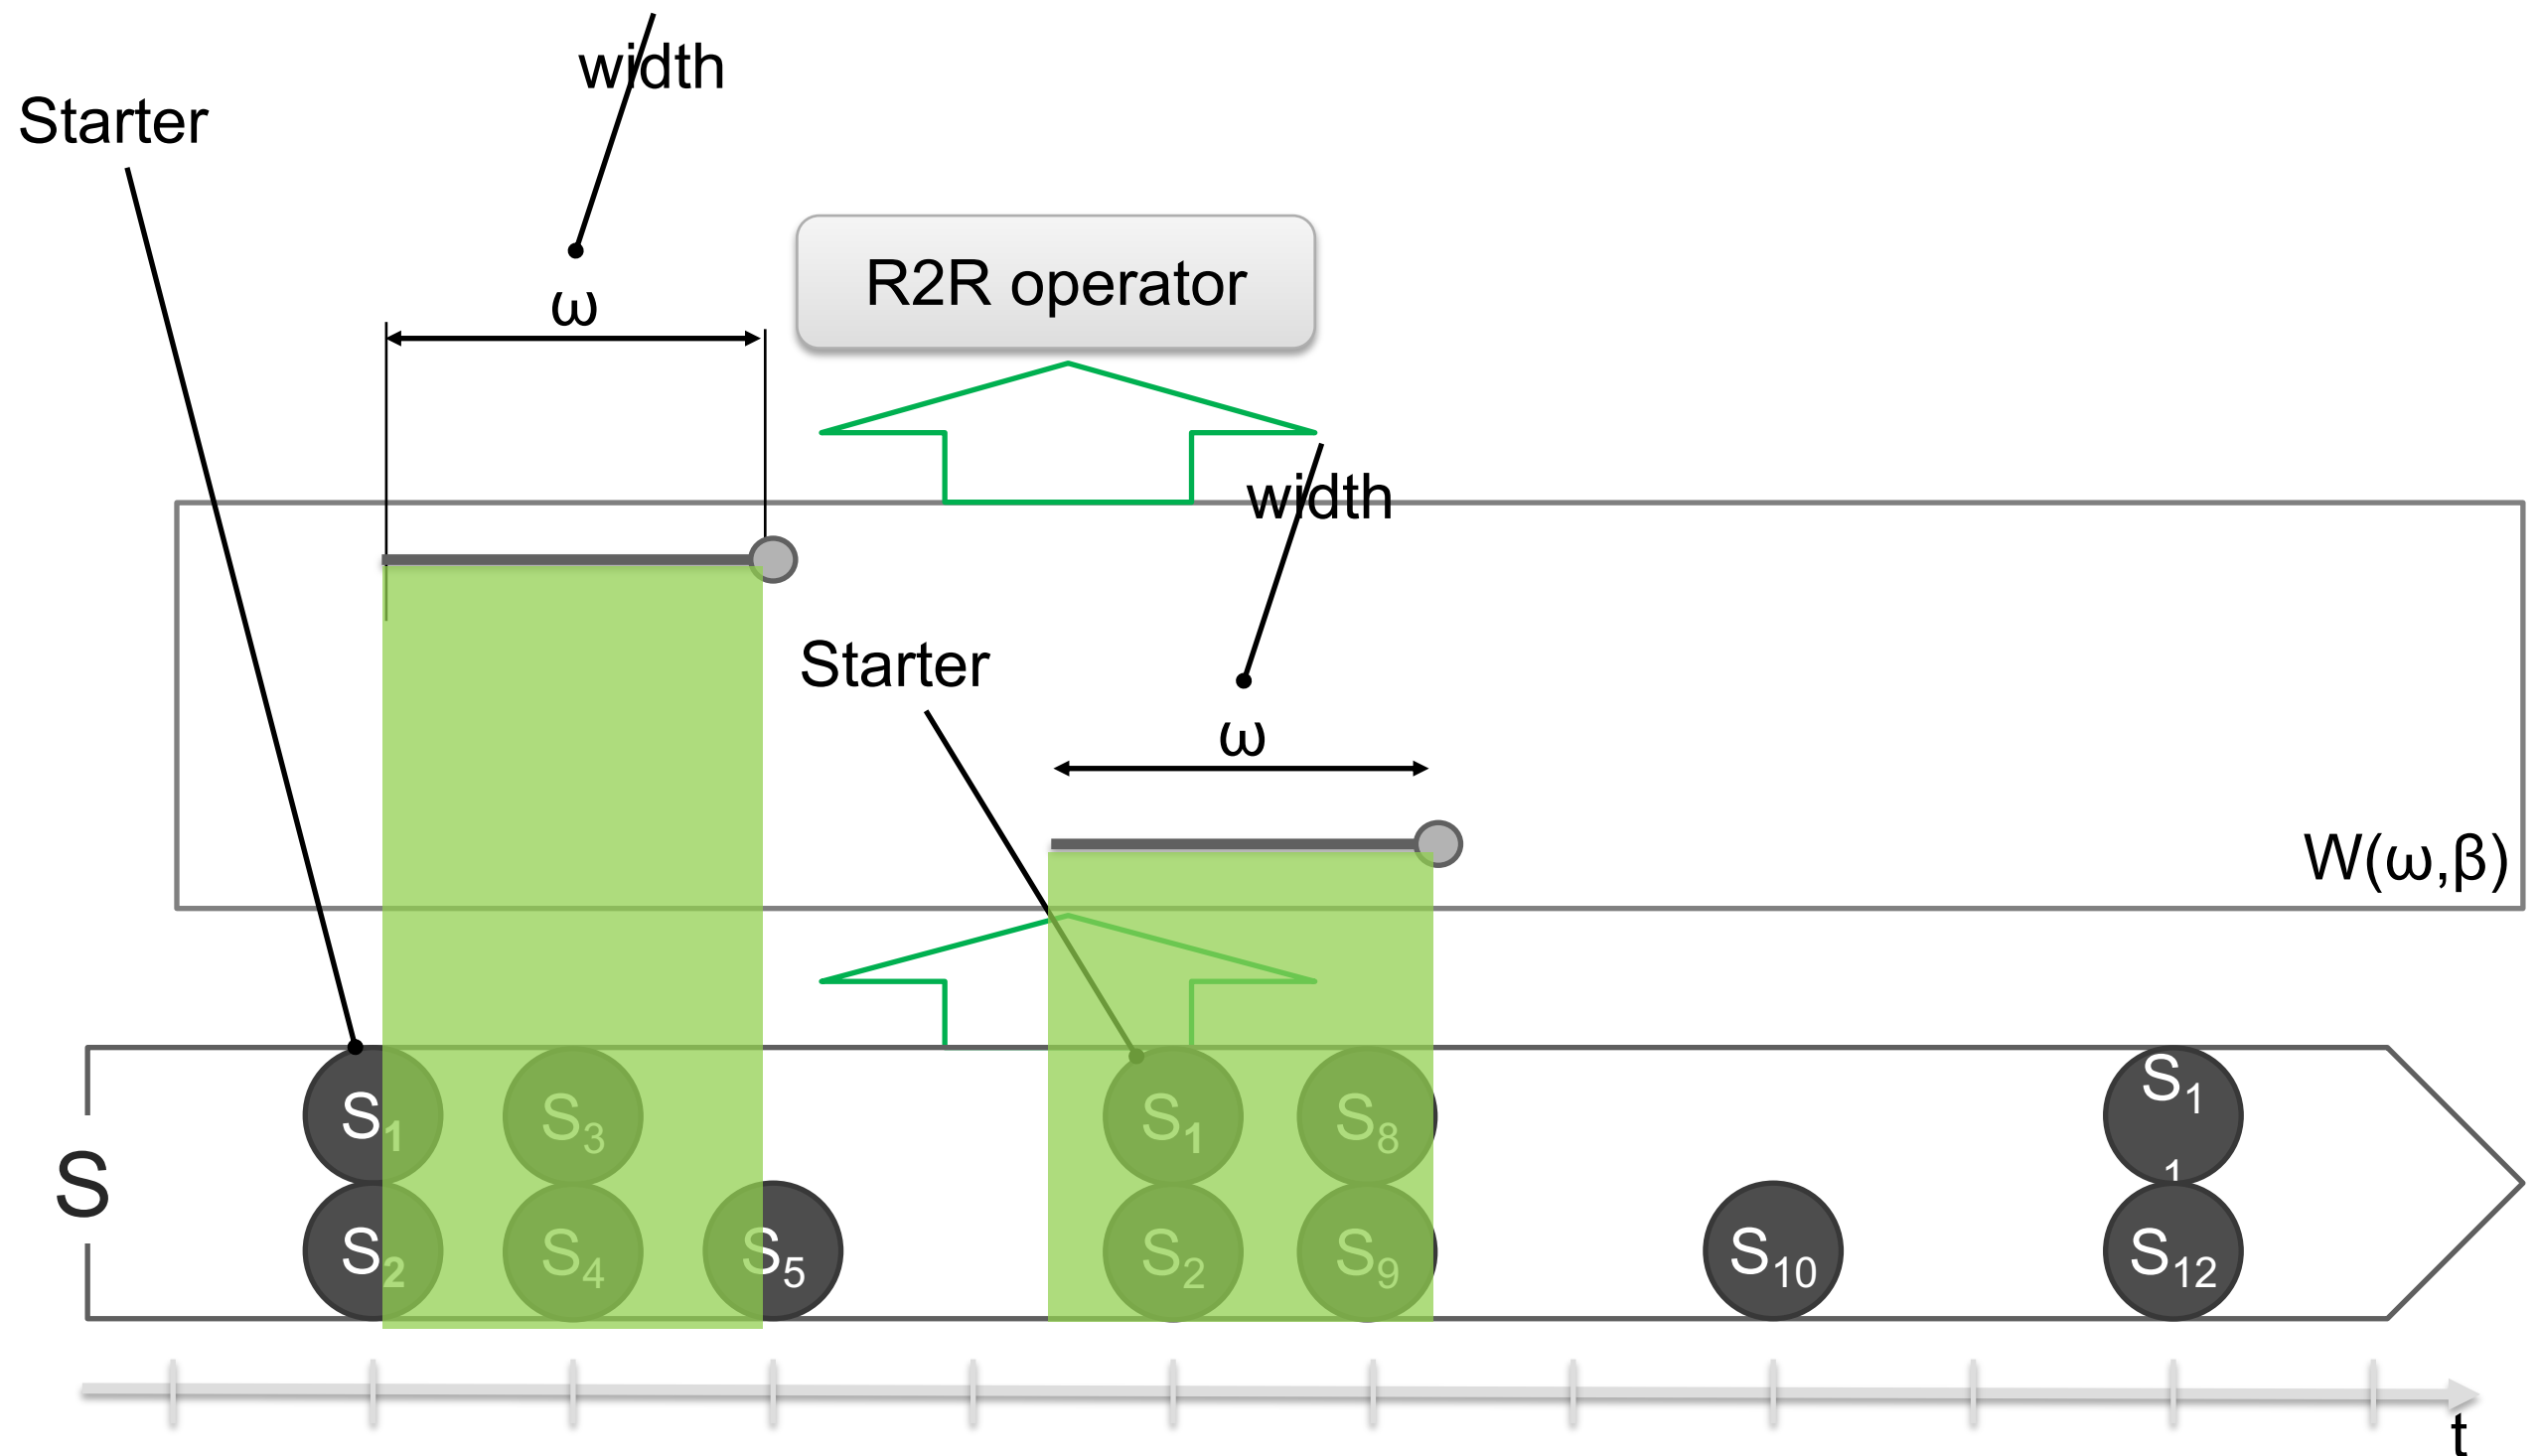
Aggregation



Tumbling Window



Session Window



KSQL Session

DDL Extension

Aggregat

```
CREATE TABLE analysis AS  
SELECT nation, COUNT (*),  
TIMESTAMPTOSTRING(windowstart(), 'yyyy-MM-dd  
HH:mm:ss') AS window_start_ts,  
TIMESTAMPTOSTRING(windowend(), 'yyyy-MM-dd  
HH:mm:ss') AS window_end_ts  
FROM pageviews WINDOW SESSION (1 MINUTE)  
GROUP BY nation;
```

Window From Function

Results

Page_82		2019-06-24 11:47:45		2019-06-24 11:47:45		1
Page_73		2019-06-24 11:47:46		2019-06-24 11:47:46		1
Page_16		2019-06-24 11:47:49		2019-06-24 11:47:49		1
Page_54		2019-06-24 11:47:25		2019-06-24 11:47:53		2
Page_68		2019-06-24 11:47:55		2019-06-24 11:47:55		1
Page_25		2019-06-24 11:47:40		2019-06-24 11:47:58		2
Page_17		2019-06-24 11:47:59		2019-06-24 11:47:59		1
Page_92		2019-06-24 11:48:02		2019-06-24 11:48:02		1
Page_83		2019-06-24 11:48:05		2019-06-24 11:48:05		1
Page_37		2019-06-24 11:48:06		2019-06-24 11:48:06		1
Page_86		2019-06-24 11:48:07		2019-06-24 11:48:07		1

Flink SQL Session

Aggregat
Custom Window Helper Functions
Group By Function

```
SELECT nation, count(*),  
SESSION_START(...), SESSION_ROWTIME(...)  
FROM pageviews  
GROUP BY SESSION(rowtime, INTERVAL 1M), nation
```

The diagram illustrates the mapping of SQL concepts to Flink SQL syntax. It features a SQL query with several components highlighted by red boxes. Arrows point from descriptive labels above to these highlighted components: 'Aggregat' points to 'count(*)', 'Custom Window Helper Functions' points to both 'SESSION_START' and 'SESSION_ROWTIME', and 'Group By Function' points to 'SESSION'.

Results

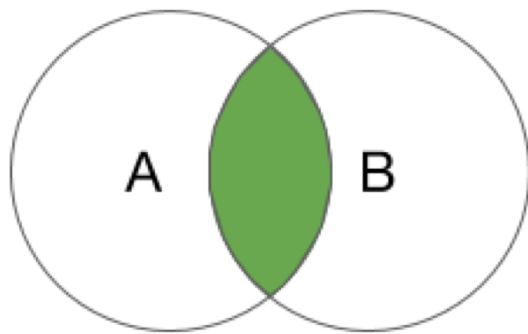
3> (Estonia,1,2019-06-24 11:52:55.538,2019-06-24 11:52:56.538,2019-06-24 11:52:56.537)
2> (Italy,1,2019-06-24 11:52:56.132,2019-06-24 11:52:57.132,2019-06-24 11:52:57.131)
1> (Egypt,1,2019-06-24 11:52:56.633,2019-06-24 11:52:57.633,2019-06-24 11:52:57.632)
3> (Estonia,1,2019-06-24 11:52:57.136,2019-06-24 11:52:58.136,2019-06-24 11:52:58.135)
2> (Italy,1,2019-06-24 11:52:57.64,2019-06-24 11:52:58.64,2019-06-24 11:52:58.639)
1> (Egypt,1,2019-06-24 11:52:58.141,2019-06-24 11:52:59.141,2019-06-24 11:52:59.14)
3> (Estonia,1,2019-06-24 11:52:58.643,2019-06-24 11:52:59.643,2019-06-24 11:52:59.642)
2> (Italy,1,2019-06-24 11:52:59.147,2019-06-24 11:53:00.147,2019-06-24 11:53:00.146)
1> (Egypt,1,2019-06-24 11:52:59.648,2019-06-24 11:53:00.648,2019-06-24 11:53:00.647)
3> (Estonia,1,2019-06-24 11:53:00.152,2019-06-24 11:53:01.152,2019-06-24 11:53:01.151)
2> (Italy,1,2019-06-24 11:53:00.653,2019-06-24 11:53:01.653,2019-06-24 11:53:01.652)
1> (Egypt,1,2019-06-24 11:53:01.158,2019-06-24 11:53:02.158,2019-06-24 11:53:02.157)

Recap

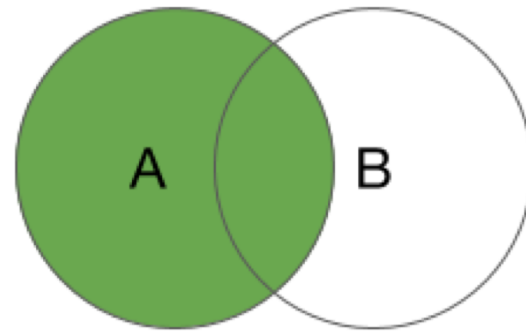
	Landmark	Tumble	Hop	Session	Aggregates
KSQL	implicit	✓	✓	✓	Standard SQL
Flink SQL	implicit	✓	✓	✓	Standard SQL
SSS	implicit	✓	X	X	Standard SQL

Table 1: Time-Based Window Operators and aggregates across different systems.

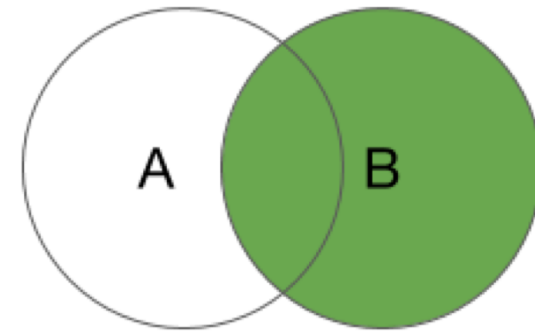
Recap on RA JOINS



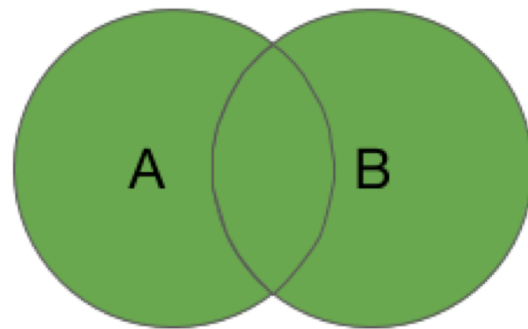
INNER JOIN



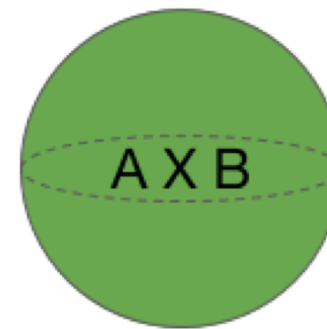
LEFT OUTER JOIN



RIGHT OUTER
JOIN



FULL OUTER
JOIN

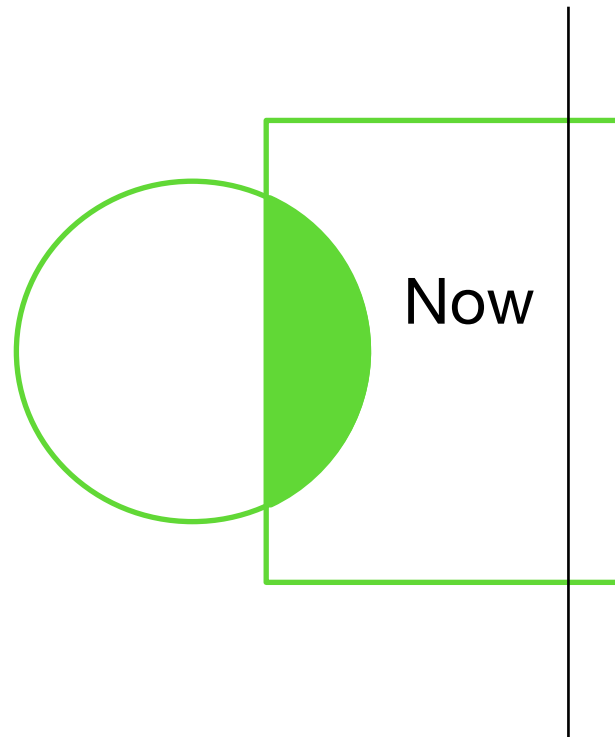


CARTESIAN
(CROSS) JOIN

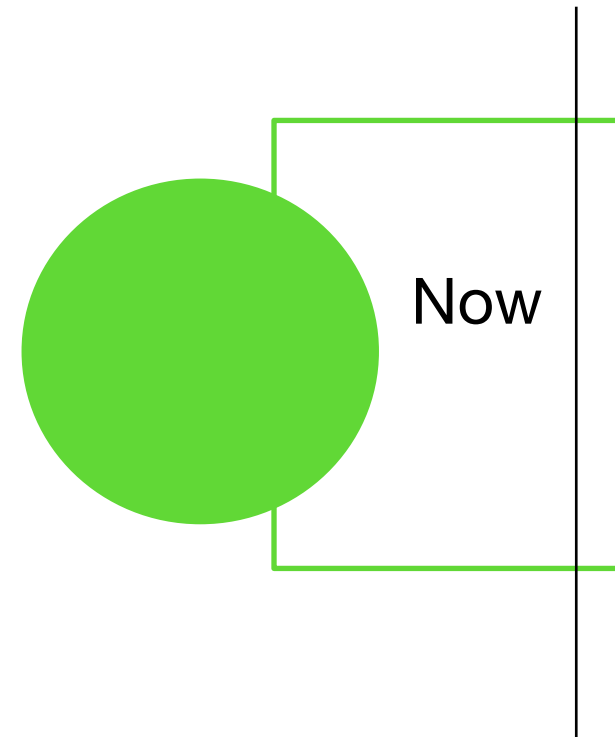
Stream-Table Joins

- Inner Joins
- Left-Outer Join
- Right-Outer Join
- Full-Outer Join

Stream-Table Joins



Inner



Left Outer

KSQL Left-Join

DDL Extension



```
CREATE STREAM SENSOR_ENRICHED AS  
SELECT S.SENSOR_ID, S.READING_VALUE, I.ITEM_ID  
FROM SENSOR_READINGS S  
LEFT JOIN ITEMS_IN_PRODUCTION I ON  
S.LINE_ID=I.LINE_ID;
```

Stream-Table Join



Flink SQL LEFT-JOIN

Stream-Table Join

```
SELECT S.SENSOR_ID, S.READING_VALUE, I.ITEM_ID  
FROM SENSOR_READINGS S LEFT JOIN  
ITEMS_IN_PRODUCTION I ON S.LINE_ID=I.LINE_ID;
```



Results

```
4> (true,0,10.12666825646483,0)
4> (true,0,10.96399203326454,0)
1> (true,2,10.874856720766067,2)
4> (true,0,10.268731915130621,0)
1> (true,2,10.786008348182463,2)
4> (true,1,10.360322470661394,1)
4> (true,0,10.809087822653261,0)
4> (true,1,10.238883138171406,1)
1> (true,2,10.776781799073452,2)
4> (true,1,10.528528144000497,1)
4> (true,0,10.532966430120872,0)
4> (true,1,10.449756056124912,1)
4> (true,1,10.66021657541424,1)
```

Spark Structured Streaming

LEFT-JOIN

Table



```
val itemsInProduction = spark.read. ...
```

Stream-Table Join



```
val sensorReadings = spark.readStream. ...
```

```
val enrichedSensorReadings =  
sensorReadings.join(itemsInProduction, "LINE_ID", "left-  
join")
```

Recap

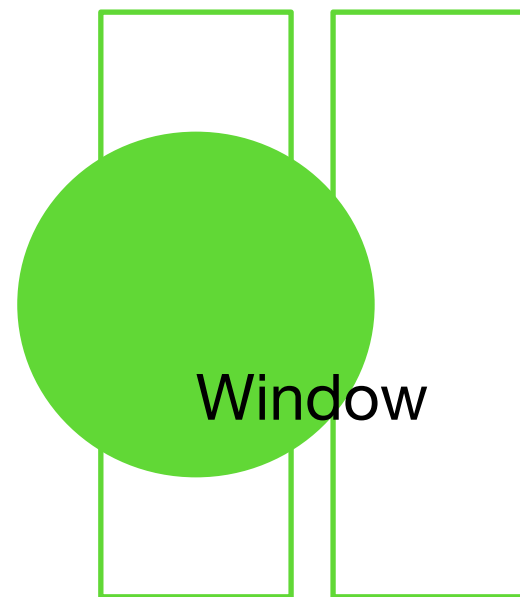
	Inner	Left Outer	Right Outer	Full Outer
KSQL	S	S	NS	NS
Flink SQL	S	S*	S*	S*
SSS	S, Stateless	S, Stateless	NS	NS

Table 2: Stream-Static Joins. [S]upported, [N]ot[S]upported. S*, Flink memory usage might grow indefinitely, Temporal Tables can be used to avoid it.

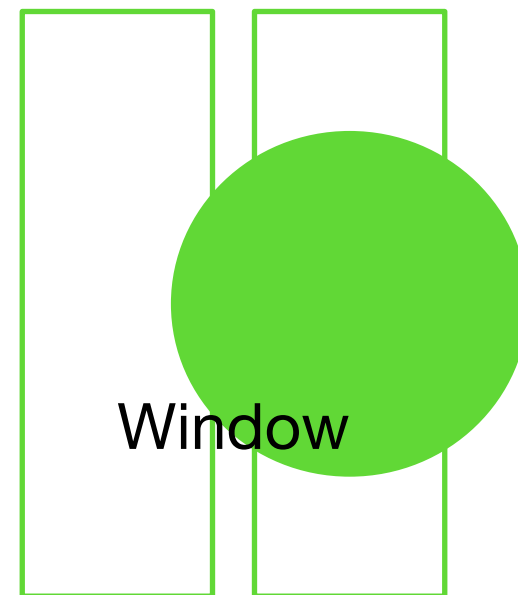
Stream-Stream Joins

- Inner Joins
- Left-Outer Join
- Right-Outer Join
- Full-Outer Join

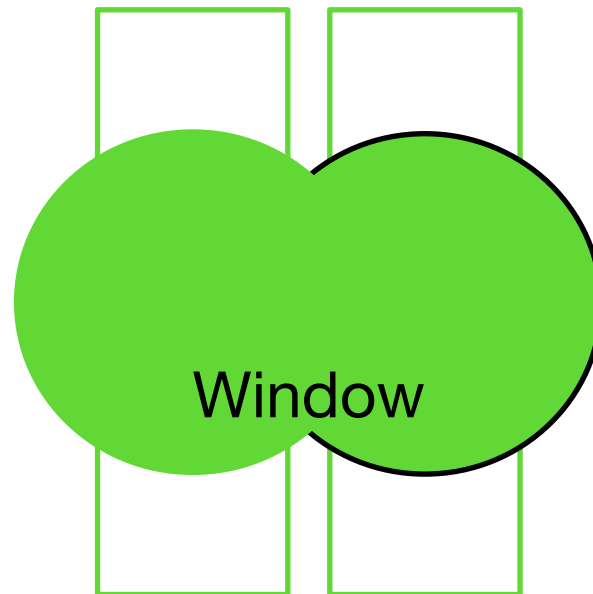
Stream-Table Joins



Left Outer



Right Outer



Full Outer

Flink SQL Inner Join

```
SELECT * FROM IMPRESSIONS, CLICKS  
WHERE IMPRESSION_ID = CLICK_ID AND  
CLICK_TIME BETWEEN IMPRESSION_TIME - INTERVAL  
'1' HOUR AND IMPRESSION_TIME
```

Spark Structured Streaming Inner Join

```
val impressions = spark.readStream. ...
```

```
val clicks = spark.readStream. ...
```

```
// Apply watermarks on event-time columns
```

```
val imprWithWtmrk  
=impressions.withWatermark("impressionTime", "2 hours")
```

```
val clicksWithWatermark =
```

```
    clicks.withWatermark("clickTime", "3 hours")
```

```
Val imprWithWtmrk.join( clicksWithWatermark, expr(""" clickAdId =  
impressionAdId AND clickTime >= impressionTime AND clickTime  
<= impressionTime + interval 1 hour"""))
```

Recap

	Inner	Left Outer	Right Outer	Full Outer
KSQL	S, win	S, win	S, win	S, win
Flink SQL	S	S*	S*	S*
SSS	S	S + w on left	S + w on right	NS

Table 3: Stream-Stream Joins. [S]upported, [N]ot[S]upported, [W]atermark, [Win]dow. S*, Flink memory usage might grow indefinitely.

Agenda

- Introduction on Stream Processing Models [done]
- Declarative Language: Opportunities, and Design Principles [done]
- Comparison of Prominent Streaming SQL Dialects for Big Stream Processing Systems [done]
- Conclusion

DEMO

KSQL and Flink

survey of spark structured streaming notebook